

Loan Approval Model

```
#importing necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# reading the data
data =
pd.read_json('/Users/yogeshdhaliya/Desktop/Task/loan_approval_dataset.
json')

df=pd.DataFrame(data)

df.head()
```

	Id	Income	Age	Experience	Married/Single	House_Ownership
0	1	1303834	23	3	single	rented
1	2	7574516	40	10	single	rented
2	3	3991815	66	4	married	rented
3	4	6256451	41	2	single	rented
4	5	5768871	47	11	single	rented

	Profession	CITY	STATE
0	Mechanical_engineer	Rewa	Madhya_Pradesh
1	Software_Developer	Parbhani	Maharashtra
2	Technical_writer	Alappuzha	Kerala
3	Software_Developer	Bhubaneswar	Odisha
4	Civil_servant	Tiruchirappalli[10]	Tamil_Nadu

	CURRENT_HOUSE_YRS	Risk_Flag
0	13	0
1	13	0
2	10	0
3	12	1
4	14	1

Data Exploration

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 252000 entries, 0 to 251999
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	Id	252000 non-null	int64
1	Income	252000 non-null	int64
2	Age	252000 non-null	int64
3	Experience	252000 non-null	int64
4	Married/Single	252000 non-null	object
5	House_Ownership	252000 non-null	object
6	Car_Ownership	252000 non-null	object
7	Profession	252000 non-null	object
8	CITY	252000 non-null	object
9	STATE	252000 non-null	object
10	CURRENT_JOB_YRS	252000 non-null	int64
11	CURRENT_HOUSE_YRS	252000 non-null	int64
12	Risk_Flag	252000 non-null	int64

```
dtypes: int64(7), object(6)
```

```
memory usage: 26.9+ MB
```

```
df.shape
```

```
(252000, 13)
```

```
df.dtypes
```

Id	int64
Income	int64
Age	int64
Experience	int64
Married/Single	object
House_Ownership	object
Car_Ownership	object
Profession	object
CITY	object
STATE	object
CURRENT_JOB_YRS	int64
CURRENT_HOUSE_YRS	int64
Risk_Flag	int64

```
dtype: object
```

```
df.isnull().sum()
```

Id	0
Income	0
Age	0

```

Experience      0
Married/Single  0
House_Ownership 0
Car_Ownership  0
Profession      0
CITY            0
STATE          0
CURRENT_JOB_YRS 0
CURRENT_HOUSE_YRS 0
Risk_Flag       0
dtype: int64

```

No missing values detected in the data

#Checking the unique values in the data

```
df.nunique()
```

```

Id      252000
Income  41920
Age      59
Experience 21
Married/Single 2
House_Ownership 3
Car_Ownership 2
Profession 51
CITY     317
STATE    29
CURRENT_JOB_YRS 15
CURRENT_HOUSE_YRS 5
Risk_Flag 2
dtype: int64

```

df.describe() #Checking the summary statistics of the data

	Id	Income	Age	Experience \
count	252000.000000	2.520000e+05	252000.000000	252000.000000
mean	126000.500000	4.997117e+06	49.954071	10.084437
std	72746.278255	2.878311e+06	17.063855	6.002590
min	1.000000	1.031000e+04	21.000000	0.000000
25%	63000.750000	2.503015e+06	35.000000	5.000000
50%	126000.500000	5.000694e+06	50.000000	10.000000
75%	189000.250000	7.477502e+06	65.000000	15.000000
max	252000.000000	9.999938e+06	79.000000	20.000000

	CURRENT_JOB_YRS	CURRENT_HOUSE_YRS	Risk_Flag
count	252000.000000	252000.000000	252000.000000
mean	6.333877	11.997794	0.123000
std	3.647053	1.399037	0.328438
min	0.000000	10.000000	0.000000
25%	3.000000	11.000000	0.000000

50%	6.000000	12.000000	0.000000
75%	9.000000	13.000000	0.000000
max	14.000000	14.000000	1.000000

Categorical Variables

```
# check how many columns are categorical
for i in df.columns:
    if df[i].dtype == 'object':
        print(i)
```

```
Married/Single
House_Ownership
Car_Ownership
Profession
CITY
STATE
```

```
df["Married/Single"].value_counts()
```

```
single    226272
married   25728
Name: Married/Single, dtype: int64
```

```
df["House_Ownership"].value_counts()
```

```
rented      231898
owned       12918
norent_noown  7184
Name: House_Ownership, dtype: int64
```

```
df["Car_Ownership"].value_counts()
```

```
no      176000
yes      76000
Name: Car_Ownership, dtype: int64
```

```
df["Profession"].value_counts()
```

```
Physician          5957
Statistician       5806
Web_designer       5397
Psychologist       5390
Computer_hardware_engineer  5372
Drafter            5359
Magistrate         5357
Fashion_Designer   5304
Air_traffic_controller  5281
Comedian           5259
Industrial_Engineer  5250
Mechanical_engineer  5217
```

Chemical_engineer	5205
Technical_writer	5195
Hotel_Manager	5178
Financial_Analyst	5167
Graphic_Designer	5166
Flight_attendant	5128
Biomedical_Engineer	5127
Secretary	5061
Software_Developer	5053
Petroleum_Engineer	5041
Police_officer	5035
Computer_operator	4990
Politician	4944
Microbiologist	4881
Technician	4864
Artist	4861
Lawyer	4818
Consultant	4808
Dentist	4782
Scientist	4781
Surgeon	4772
Aviator	4758
Technology_specialist	4737
Design_Engineer	4729
Surveyor	4714
Geologist	4672
Analyst	4668
Army_officer	4661
Architect	4657
Chef	4635
Librarian	4628
Civil_engineer	4616
Designer	4598
Economist	4573
Firefighter	4507
Chartered_Accountant	4493
Civil_servant	4413
Official	4087
Engineer	4048

Name: Profession, dtype: int64

df["CITY"].value_counts()

Vijayanagaram	1259
Bhopal	1208
Bulandshahr	1185
Saharsa[29]	1180
Vijayawada	1172
...	
Ujjain	486

```

Warangal[11][12]      459
Bettiah[33]           457
Katni                  448
Karaikudi              431
Name: CITY, Length: 317, dtype: int64

```

```
df["STATE"].value_counts()
```

```

Uttar_Pradesh      28400
Maharashtra        25562
Andhra_Pradesh     25297
West_Bengal        23483
Bihar              19780
Tamil_Nadu         16537
Madhya_Pradesh     14122
Karnataka          11855
Gujarat            11408
Rajasthan          9174
Jharkhand          8965
Haryana            7890
Telangana          7524
Assam              7062
Kerala             5805
Delhi              5490
Punjab             4720
Odisha             4658
Chhattisgarh       3834
Uttarakhand        1874
Jammu_and_Kashmir  1780
Puducherry         1433
Mizoram            849
Manipur            849
Himachal_Pradesh   833
Tripura            809
Uttar_Pradesh[5]   743
Chandigarh         656
Sikkim             608
Name: STATE, dtype: int64

```

Feature Selection - chi square test - to check association between categorical features and target feature

```

from scipy.stats import chi2_contingency
for i in
['Married/Single', 'House_Ownership', 'Car_Ownership', 'Profession', 'CITY',
', 'STATE']:
    chi2, pval, _, _ = chi2_contingency(pd.crosstab(df[i],
df['Risk_Flag']))
    print(i, '---', pval)

```

```
Married/Single --- 3.773053705715196e-26
House_Ownership --- 1.8381930028370595e-40
Car_Ownership --- 1.7350853850183746e-33
Profession --- 5.108641602000937e-98
CITY --- 0.0
STATE --- 2.0057472384130266e-136
```

Since all the categorical feature have p-value ≤ 0.05 , we will accept all

Numerical Variables

```
# check how many columns are numerical
for i in df.columns:
    if df[i].dtype == 'int':
        print(i)

Id
Income
Age
Experience
CURRENT_JOB_YRS
CURRENT_HOUSE_YRS
Risk_Flag

from scipy.stats import pointbiserialr

# Calculate point biserial correlation for each integer variable
corr1, p_value1 = pointbiserialr(df['Income'], df['Risk_Flag'])
corr2, p_value2 = pointbiserialr(df['Age'], df['Risk_Flag'])
corr3, p_value3 = pointbiserialr(df['Experience'], df['Risk_Flag'])
corr4, p_value4 = pointbiserialr(df['CURRENT_JOB_YRS'],
df['Risk_Flag'])
corr5, p_value5 = pointbiserialr(df['CURRENT_HOUSE_YRS'],
df['Risk_Flag'])

print(f'Correlation between integer_variable1 and Target: {corr1}, p-
value: {p_value1}')
print(f'Correlation between integer_variable2 and Target: {corr2}, p-
value: {p_value2}')
print(f'Correlation between integer_variable2 and Target: {corr3}, p-
value: {p_value3}')
print(f'Correlation between integer_variable2 and Target: {corr4}, p-
value: {p_value4}')
print(f'Correlation between integer_variable2 and Target: {corr5}, p-
value: {p_value5}')
```

```
Correlation between integer_variable1 and Target: -
0.003091168122271239, p-value: 0.12072203190827889
Correlation between integer_variable2 and Target: -
0.02180927605070088, p-value: 6.687623645092961e-28
```

```
Correlation between integer_variable2 and Target: -
0.03452261289070651, p-value: 2.5475634857843415e-67
Correlation between integer_variable2 and Target: -
0.01694158205212183, p-value: 1.8132212225854788e-17
Correlation between integer_variable2 and Target: -
0.0043751630776717836, p-value: 0.028069472763846918
```

Interpretation of Results:

1. Income:
 - Correlation: -0.0031
 - p-value: 0.1207
 - Interpretation: Weak correlation and not statistically significant ($p > 0.05$).
1. Age:
 - Correlation: -0.0218
 - p-value: 6.6876e-28
 - Interpretation: Weak correlation but statistically significant ($p < 0.05$).
1. Experience:
 - Correlation: -0.0345
 - p-value: 2.5476e-67
 - Interpretation: Weak correlation but statistically significant ($p < 0.05$).
1. CURRENT_JOB_YRS:
 - Correlation: -0.0169
 - p-value: 1.8132e-17
 - Interpretation: Weak correlation but statistically significant ($p < 0.05$).
1. CURRENT_HOUSE_YRS:
 - Correlation: -0.0044
 - p-value: 0.0281
 - Interpretation: Very weak correlation but statistically significant ($p < 0.05$).

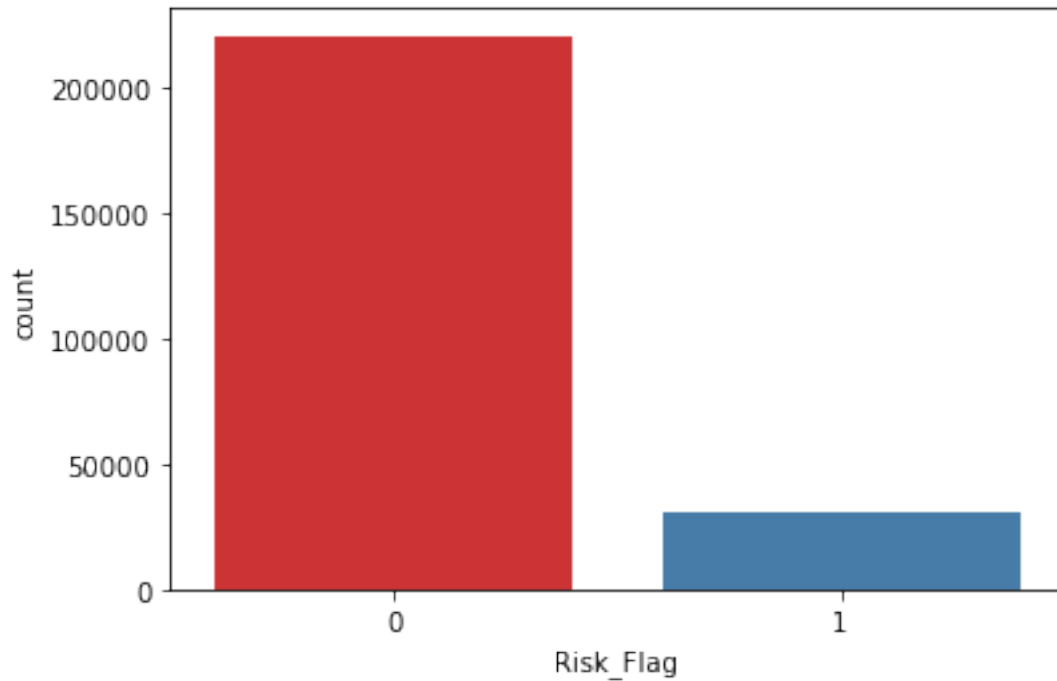
Data Visualization

```
# checking distribution between target variables
sns.countplot(df['Risk_Flag'], palette='Set1')
```

```
/opt/anaconda3/lib/python3.9/site-packages/seaborn/_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From
version 0.12, the only valid positional argument will be `data`, and
passing other arguments without an explicit keyword will result in an
error or misinterpretation.
```

```
warnings.warn(
```

```
<AxesSubplot:xlabel='Risk_Flag', ylabel='count'>
```





```
sns.countplot(df['Married/Single'],palette='Set1')
```

```
/opt/anaconda3/lib/python3.9/site-packages/seaborn/_decorators.py:36:  
FutureWarning: Pass the following variable as a keyword arg: x. From  
version 0.12, the only valid positional argument will be `data`, and  
passing other arguments without an explicit keyword will result in an  
error or misinterpretation.  
warnings.warn(  

```

```
<AxesSubplot:xlabel='Married/Single', ylabel='count'>
```




```
sns.countplot(df['House_Ownership'],palette='Set1')
```

```
/opt/anaconda3/lib/python3.9/site-packages/seaborn/_decorators.py:36:  
FutureWarning: Pass the following variable as a keyword arg: x. From  
version 0.12, the only valid positional argument will be `data`, and  
passing other arguments without an explicit keyword will result in an  
error or misinterpretation.
```

```
warnings.warn(
```

```
<AxesSubplot:xlabel='House_Ownership', ylabel='count'>
```




```
sns.countplot(df['Car_Ownership'],palette='Set1')
```

```
/opt/anaconda3/lib/python3.9/site-packages/seaborn/_decorators.py:36:  
FutureWarning: Pass the following variable as a keyword arg: x. From  
version 0.12, the only valid positional argument will be `data`, and  
passing other arguments without an explicit keyword will result in an  
error or misinterpretation.
```

```
warnings.warn(
```


```
<AxesSubplot:xlabel='Car_Ownership', ylabel='count'>
```



```
plt.figure(figsize=(12, 12))
ax = sns.countplot(y=df['STATE'])

plt.title("Profession Frequency")
plt.ylabel("Profession Type")
plt.xlabel("Count")


for container in ax.containers:
    ax.bar_label(container)
plt.show()
```



```
plt.figure(figsize=(12, 20))
ax = sns.countplot(y=df['Profession'])

plt.title("Profession Frequency")
plt.ylabel("Profession Type")
plt.xlabel("Count")

for container in ax.containers:
    ax.bar_label(container)
plt.show()
```

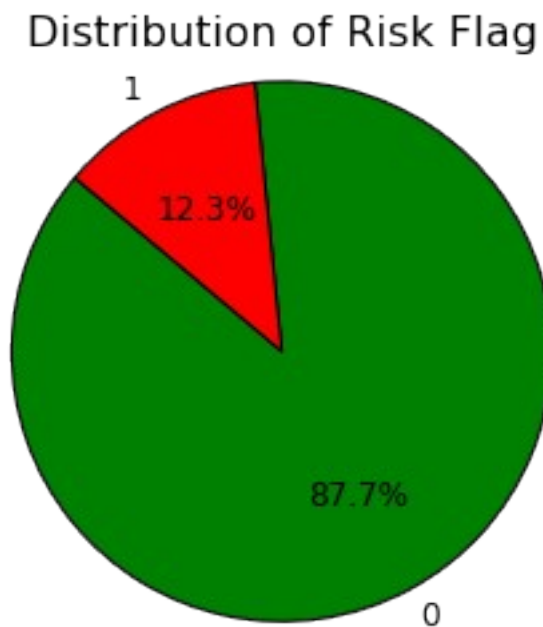


```
plt.figure(figsize=(4, 4))

counts = df['Risk_Flag'].value_counts().sort_values(ascending=False)
custom_colors = ['green', 'red']

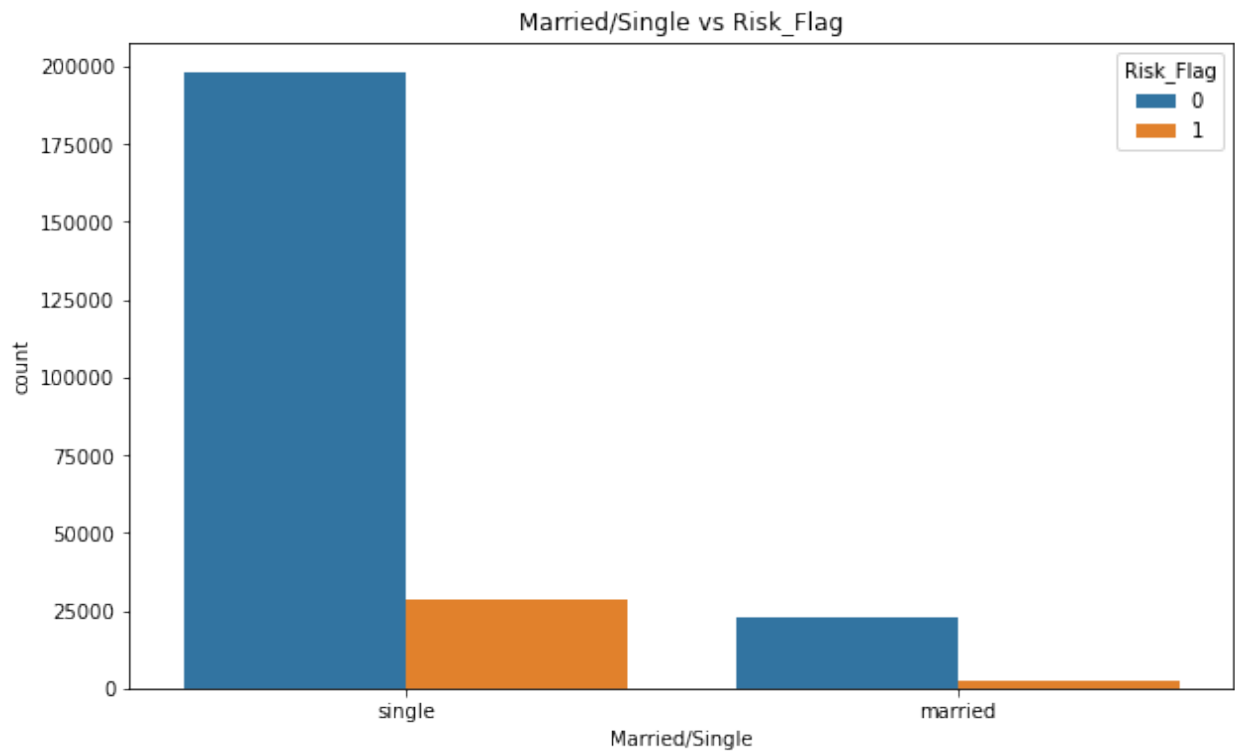
plt.pie(counts, labels=counts.index, autopct='%1.1f%%',
        startangle=140, colors=custom_colors,
        wedgeprops={'edgecolor': 'black', 'linewidth': 1},
        textprops={'fontsize': 12})

plt.axis('equal')
plt.title('Distribution of Risk Flag', fontsize=16)
plt.show()
```




In target variable 12.3% of clients are high risk individuals & 87.7% of clients are low risk individuals.

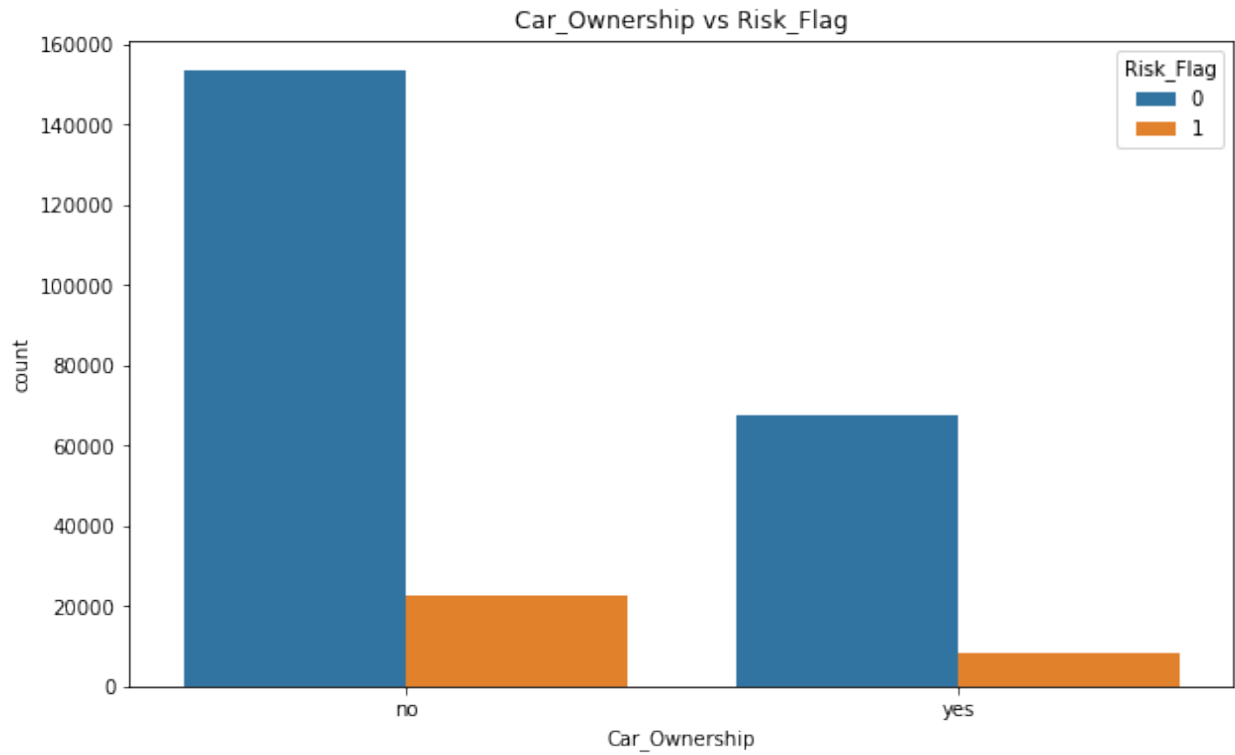
```
plt.figure(figsize=(10, 6))
sns.countplot(x='Married/Single', hue='Risk_Flag', data=df)
plt.title('Married/Single vs Risk_Flag')
plt.show()
```




```
plt.figure(figsize=(10, 6))
sns.countplot(x='House_Ownership', hue='Risk_Flag', data=df)
plt.title('House_Ownership vs Risk_Flag')
plt.show()
```


```
plt.figure(figsize=(10, 6))
sns.countplot(x='Car_Ownership', hue='Risk_Flag', data=df)
plt.title('Car_Ownership vs Risk_Flag')
plt.show()
```




```
# Histogram of Income  
plt.figure(figsize=(10, 6))  
sns.histplot(df['Income'], bins=30, kde=True)  
plt.title('Distribution of Income')  
plt.show()
```



```
# Histogram of Age
plt.figure(figsize=(10, 6))
sns.histplot(df['Age'], bins=30, kde=True)
plt.title('Distribution of Age')
plt.show()
```



```
plt.figure(figsize=(8,8))
sns.heatmap(df.corr(),annot=True,cmap='coolwarm')
<AxesSubplot:>
```



```
#Splitting the data into X and y
from sklearn.model_selection import train_test_split
X=df.drop('Risk_Flag',axis=1)
y=df['Risk_Flag']
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=42)
X_val,X_test,y_val,y_test=train_test_split(X_test,y_test,test_size=0.5,random_state=42)

from sklearn.preprocessing import LabelEncoder
```

```

# Initialize LabelEncoder
le = LabelEncoder()

categorical_columns = X.select_dtypes(include=['object']).columns

# Fit and transform categorical columns in training set
for col in categorical_columns:
    X_train[col] = le.fit_transform(X_train[col])

# Transform categorical columns in validation set, handle unknown labels
for col in categorical_columns:
    # Handle unknown labels by assigning them a default value (e.g., 0)
    X_val[col] = le.transform(X_val[col].map(lambda x: 0 if x not in le.classes_ else x))

# Transform categorical columns in testing set, handle unknown labels
for col in categorical_columns:
    # Handle unknown labels by assigning them a default value (e.g., 0)
    X_test[col] = le.transform(X_test[col].map(lambda x: 0 if x not in le.classes_ else x))

# Print the encoded datasets
print("X_train:")
print(X_train.head())

print("\nX_val:")
print(X_val.head())

print("\nX_test:")
print(X_test.head())

```

```

-----
-----
TypeError                                Traceback (most recent call
last)
Input In [35], in <cell line: 13>()
    12 # Transform categorical columns in validation set, handle
unknown labels
    13 for col in categorical_columns:
    14     # Handle unknown labels by assigning them a default value
(e.g., 0)
--> 15     X_val[col] = le.transform(X_val[col].map(lambda x: 0 if x
not in le.classes_ else x))
    17 # Transform categorical columns in testing set, handle unknown
labels
    18 for col in categorical_columns:
    19     # Handle unknown labels by assigning them a default value

```

(e.g., 0)

```
File
/opt/anaconda3/lib/python3.9/site-packages/sklearn/preprocessing/_label.py:138, in LabelEncoder.transform(self, y)
    135 if _num_samples(y) == 0:
    136     return np.array([])
--> 138 return _encode(y, uniques=self.classes_)
```

```
File
/opt/anaconda3/lib/python3.9/site-packages/sklearn/utils/_encode.py:187, in _encode(values, uniques, check_unknown)
    185 else:
    186     if check_unknown:
--> 187         diff = _check_unknown(values, uniques)
    188         if diff:
    189             raise ValueError(f"y contains previously unseen
labels: {str(diff)}")
```

```
File
/opt/anaconda3/lib/python3.9/site-packages/sklearn/utils/_encode.py:261, in _check_unknown(values, known_values, return_mask)
    258     valid_mask = np.ones(len(values), dtype=bool)
    260 # check for nans in the known_values
--> 261 if np.isnan(known_values).any():
    262     diff_is_nan = np.isnan(diff)
    263     if diff_is_nan.any():
    264         # removes nan from valid_mask
```

TypeError: ufunc 'isnan' not supported for the input types, and the inputs could not be safely coerced to any supported types according to the casting rule 'safe'

```
from sklearn.preprocessing import StandardScaler

numerical_features =
X.select_dtypes(include=['int64', 'float64']).columns

# Initialize StandardScaler
scaler = StandardScaler()

# Scale numerical features in training set
X_train[numerical_features] =
scaler.fit_transform(X_train[numerical_features])

# Scale numerical features in validation set
X_val[numerical_features] =
scaler.transform(X_val[numerical_features])

# Scale numerical features in testing set
```

```
X_test[numerical_features] =
scaler.transform(X_test[numerical_features])
```

```
X_train.head()
```

	Id	Income	Age	Experience	Married/Single	\
200471	1.024787	1.430422	1.170436	1.318759		1
92611	-0.457717	-0.820701	-0.644922	-1.680840		1
86397	-0.543127	-1.385148	-0.644922	-0.181041		1
110500	-0.211838	-0.383000	-1.699001	-1.347551		1
185133	0.813971	1.551521	-0.527803	-0.847618		1

	House_Ownership	Car_Ownership	Profession	CITY	STATE	\
200471	2	0	4	270	20	
92611	2	0	13	311	22	
86397	2	0	42	123	23	
110500	2	0	48	295	1	
185133	0	1	50	56	11	

	CURRENT_JOB_YRS	CURRENT_HOUSE_YRS
200471	-0.914084	-0.712569
92611	-1.736837	-0.712569
86397	-0.639833	0.717243
110500	-1.188335	0.002337
185133	-0.365582	0.717243

```
# Remove the 'Id' column from all datasets
```

```
X_train = X_train.drop('Id', axis=1)
```

```
X_val = X_val.drop('Id', axis=1)
```

```
X_test = X_test.drop('Id', axis=1)
```

```
#Saving the preprocessed data
```

```
X_train.to_csv('/Users/yogeshdhaliya/Desktop/Task/X_train.csv',index=False)
```

```
X_val.to_csv('/Users/yogeshdhaliya/Desktop/Task/X_val.csv',index=False)
```

```
X_test.to_csv('/Users/yogeshdhaliya/Desktop/Task/X_test.csv',index=False)
```

```
y_train.to_csv('/Users/yogeshdhaliya/Desktop/Task/y_train.csv',index=False)
```

```
y_val.to_csv('/Users/yogeshdhaliya/Desktop/Task/y_val.csv',index=False)
```

```
y_test.to_csv('/Users/yogeshdhaliya/Desktop/Task/y_test.csv',index=False)
```

```
import pandas as pd
```

```
# Load the preprocessed data
```

```
X_train = pd.read_csv("/Users/yogeshdhaliya/Desktop/Task/X_train.csv")
```

```
X_val = pd.read_csv('/Users/yogeshdhaliya/Desktop/Task/X_val.csv')
```

```
X_test = pd.read_csv('/Users/yogeshdhaliya/Desktop/Task/X_test.csv')
```



```

y_train = pd.read_csv('/Users/yogeshdhaliya/Desktop/Task/y_train.csv')
y_val = pd.read_csv('/Users/yogeshdhaliya/Desktop/Task/y_val.csv')
y_test = pd.read_csv('/Users/yogeshdhaliya/Desktop/Task/y_test.csv')

# Determine the number of classes in the target variable
n_classes = len(y_train['Risk_Flag'].unique())
print(f'Number of classes: {n_classes}')

Number of classes: 2

#maximum number of LDA components(Applying LDA)
max_components = n_classes - 1
print(f'Maximum number of LDA components: {max_components}')

Maximum number of LDA components: 1

from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_val_scaled = scaler.transform(X_val)
X_test_scaled = scaler.transform(X_test)

# Apply LDA
lda = LinearDiscriminantAnalysis(n_components=max_components)
# Fit LDA and transform the training data
X_train_lda = lda.fit_transform(X_train_scaled,
y_train.values.ravel())
X_val_lda = lda.transform(X_val_scaled)
X_test_lda = lda.transform(X_test_scaled)

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
# Convert the transformed arrays back to DataFrames for easier
inspection
X_train_lda_df = pd.DataFrame(X_train_lda, columns=[f'LDA{i+1}' for i
in range(X_train_lda.shape[1])])
X_val_lda_df = pd.DataFrame(X_val_lda, columns=[f'LDA{i+1}' for i in
range(X_val_lda.shape[1])])
X_test_lda_df = pd.DataFrame(X_test_lda, columns=[f'LDA{i+1}' for i in
range(X_test_lda.shape[1])])
# Train and evaluate a classifier using the LDA-transformed data
clf = RandomForestClassifier(random_state=42)
clf.fit(X_train_lda_df, y_train.values.ravel())
y_pred_val = clf.predict(X_val_lda_df)
y_pred_test = clf.predict(X_test_lda_df)
accuracy = accuracy_score(y_val, y_pred_val)
accuracy1=accuracy_score(y_test,y_pred_test)
print(f'Validation Accuracy with LDA: {accuracy}')
print(f'Test Accuracy with LDA: {accuracy1}')

```

Validation Accuracy with LDA: 0.720515873015873

Test Accuracy with LDA: 0.7209126984126984

#using smote to balance the data

```
from imblearn.over_sampling import SMOTE
```

```
smote = SMOTE(random_state=42)
```

```
X_train_smote, y_train_smote = smote.fit_resample(X_train_lda_df,  
y_train)
```

Convert the transformed arrays back to DataFrames for easier inspection

```
X_train_smote = pd.DataFrame(X_train_smote, columns=[f'LDA{i+1}' for i  
in range(X_train_smote.shape[1])])
```

#Random forest with just train_scaled data

```
clf = RandomForestClassifier(random_state=42)
```

```
clf.fit(X_train_scaled, y_train.values.ravel())
```

```
y_pred_val = clf.predict(X_val_scaled)
```

```
accuracy = accuracy_score(y_val, y_pred_val)
```

```
print(f'Validation Accuracy with just scaled data: {accuracy}')
```

Validation Accuracy with just scaled data: 0.8740079365079365

#using XGBoost

```
from xgboost import XGBClassifier
```

```
clf = XGBClassifier(random_state=42)
```

```
clf.fit(X_train_scaled, y_train.values.ravel())
```

```
y_pred_val = clf.predict(X_val_scaled)
```

```
accuracy = accuracy_score(y_val, y_pred_val)
```

```
y_pred_test = clf.predict(X_test_scaled)
```

```
accuracy1=accuracy_score(y_test,y_pred_test)
```

```
print(f'Validation Accuracy with XGBoost: {accuracy}')
```

```
print(f'Test Accuracy with XGBoost: {accuracy1}')
```

Validation Accuracy with XGBoost: 0.8649206349206349

Test Accuracy with XGBoost: 0.8686111111111111

#Using XGBoost with LDA

```
clf = XGBClassifier(random_state=42)
```

```
clf.fit(X_train_lda_df, y_train.values.ravel())
```

```
y_pred_val = clf.predict(X_val_lda_df)
```

```
accuracy = accuracy_score(y_val, y_pred_val)
```

```
y_pred_test = clf.predict(X_test_lda_df)
```

```
accuracy1=accuracy_score(y_test,y_pred_test)
```

```
print(f'Validation Accuracy with XGBoost and LDA: {accuracy}')
```

```
print(f'Test Accuracy with XGBoost and LDA: {accuracy1}')
```

Validation Accuracy with XGBoost and LDA: 0.8703571428571428

Test Accuracy with XGBoost and LDA: 0.8740476190476191

#Using logistic regression

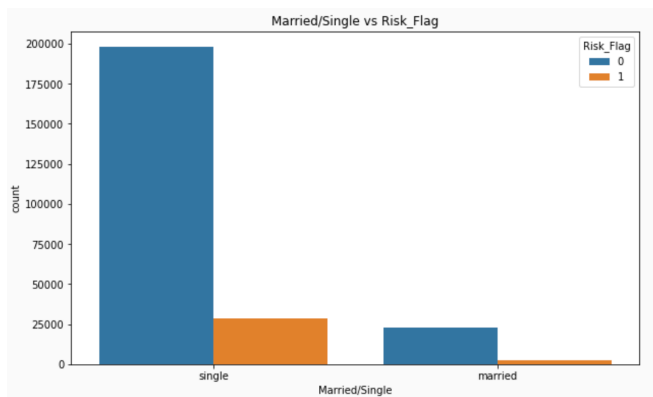
```
from sklearn.linear_model import LogisticRegression
```

```
clf = LogisticRegression(random_state=42)
clf.fit(X_train_scaled, y_train.values.ravel())
y_pred_val = clf.predict(X_val_scaled)
accuracy = accuracy_score(y_val, y_pred_val)
y_pred_test = clf.predict(X_test_scaled)
accuracy1=accuracy_score(y_test,y_pred_test)
print(f'Validation Accuracy with Logistic Regression: {accuracy}')
print(f'Test Accuracy with Logistic Regression: {accuracy1}')
```

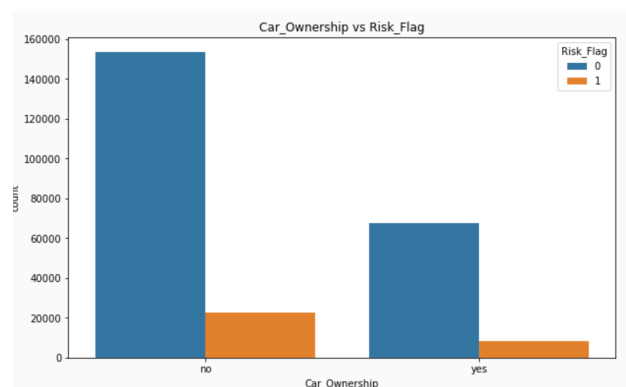
Validation Accuracy with Logistic Regression: 0.8740079365079365
Test Accuracy with Logistic Regression: 0.8778571428571429

After trying out various ML models Random Forest Classifier & Logistic Regression got the highest accuracy result. I have got accuracy score of 87%

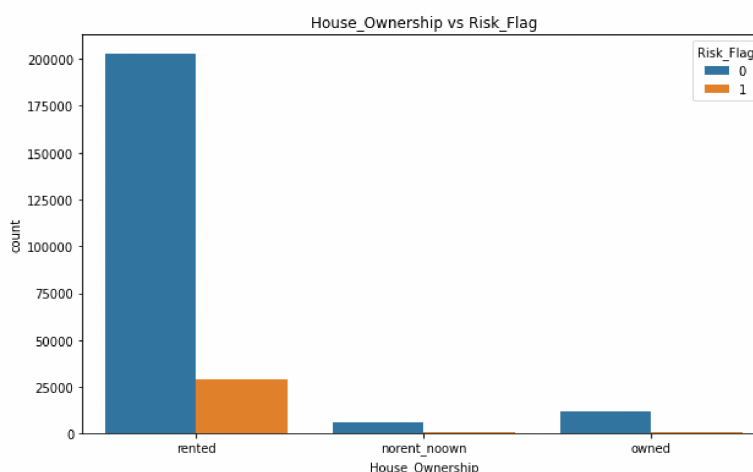
What are the main deciding factors associated with risk?



Single applicants exhibit a higher risk profile compared to married ones. This aligns with real-world observations, as factors like shared finances and stability often come with marriage.



Car ownership is associated with a lower risk of loan default. This suggests potential borrowers with cars might have better financial stability or a stronger ability to repay loans.



Individuals renting a house exhibit a higher risk profile compared to homeowners. This aligns with the potential for greater financial stability and asset ownership associated with homeownership.

In conclusion, our analysis reveals several key factors influencing loan risk. Marital status, car ownership, and homeownership all play a significant role. Single applicants, individuals without car ownership, and renters tend to be associated with higher risk profiles. These findings likely reflect factors like shared finances, financial stability, and asset ownership, which can influence a borrower's ability to repay loans.

Report By:

Yogesh Dhaliya
6350214104
yogeshdhaliyaa@gmail.com
<http://linkedin.com/in/yogeshdhaliyaa>
<https://yogeshdhaliya.github.io/yogeshportfolio.github.io/>